

Assignment 1

Machine Learning 1, SS24

Team Members		
Last name	First name	Matriculation Number
Farfan Roca	Javier	12130995
Omanovic	Melvis	1600002

1 Linear Regression – Detection of memristor faults

$$\textcircled{1} E(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(\Delta R_i^{\text{ideal}}) - \Delta R_i)^2$$

$$\theta^* = \arg \min_{\theta} E(\theta)$$

$$\frac{\partial E(\theta)}{\partial \theta} = \frac{1}{m} \sum_{i=1}^m 2 \cdot (\theta \Delta R_i^{\text{ideal}} - \Delta R_i) \cdot \Delta R_i^{\text{ideal}}$$

$$\frac{1}{m} \sum_{i=1}^m 2 \cdot (\theta^* \Delta R_i^{\text{ideal}} - \Delta R_i) \cdot \Delta R_i^{\text{ideal}} = 0 \quad / \cdot \frac{m}{2}$$

$$\theta^* \cdot \sum_{i=1}^m (\Delta R_i^{\text{ideal}})^2 - \sum_{i=1}^m \Delta R_i \cdot \Delta R_i^{\text{ideal}} = 0$$

$$\theta^* \cdot \sum_{i=1}^m (\Delta R_i^{\text{ideal}})^2 = \sum_{i=1}^m \Delta R_i \cdot \Delta R_i^{\text{ideal}} \quad / : \sum_{i=1}^m (\Delta R_i^{\text{ideal}})^2$$

$$\theta^* = \frac{\sum_{i=1}^m \Delta R_i \cdot \Delta R_i^{\text{ideal}}}{\sum_{i=1}^m (\Delta R_i^{\text{ideal}})^2}$$

$$\textcircled{2} E(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta_0, \theta_1}(\Delta R_i^{\text{ideal}}) - \Delta R_i)^2,$$

$$(\theta_0^*, \theta_1^*) = \arg \min_{\theta_0, \theta_1} E(\theta_0, \theta_1)$$

$$\cancel{E(\theta_0, \theta_1)} E(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (\theta_0 + \theta_1 \Delta R_i^{\text{ideal}} - \Delta R_i)^2$$

$$\frac{\partial E(\theta_0, \theta_1)}{\partial \theta_0} = \frac{1}{m} \sum_{i=1}^m 2 \cdot (\theta_0 + \theta_1 \Delta R_i^{\text{ideal}} - \Delta R_i) \cdot 1$$

$$= \frac{2}{m} \sum_{i=1}^m (\theta_0 + \theta_1 \Delta R_i^{\text{ideal}} - \Delta R_i)$$

$$\frac{\partial E(\theta_0, \theta_1)}{\partial \theta_1} = \frac{1}{m} \sum_{i=1}^m 2 \cdot (\theta_0 + \theta_1 \Delta R_i^{\text{ideal}} - \Delta R_i) \cdot \Delta R_i^{\text{ideal}}$$

$$= \frac{2}{m} \sum_{i=1}^m (\theta_0 + \theta_1 \Delta R_i^{\text{ideal}} - \Delta R_i) \cdot \Delta R_i^{\text{ideal}}$$

$$\frac{2}{m} \sum_{i=1}^m (\theta_0^* + \theta_1^* \cdot R_i^{\text{ideal}} - \Delta R_i) = 0$$

$$m \cdot \theta_0^* + \theta_1^* \sum_{i=1}^m \Delta R_i^{\text{ideal}} - \sum_{i=1}^m \Delta R_i = 0$$

$$m \cdot \theta_0^* = \sum_{i=1}^m \Delta R_i - \theta_1^* \sum_{i=1}^m \Delta R_i^{\text{ideal}}$$

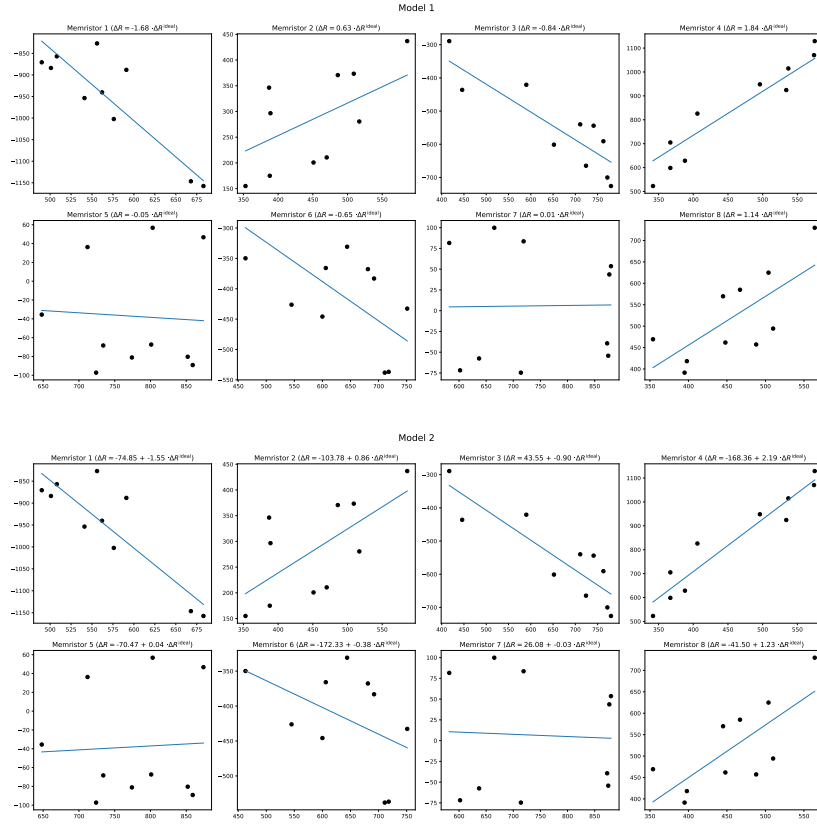
$$\theta_0^* = \frac{\sum_{i=1}^m \Delta R_i - \theta_1^* \sum_{i=1}^m \Delta R_i^{\text{ideal}}}{m}$$

$$\frac{2}{m} \sum_{i=1}^m (\theta_0^* + \theta_1^* \cdot \Delta R_i^{\text{ideal}} - \Delta R_i) \cdot \Delta R_i^{\text{ideal}} = 0$$

$$\theta_0^* \cdot \sum_{i=1}^m \Delta R_i^{\text{ideal}} + \theta_1^* \sum_{i=1}^m (\Delta R_i^{\text{ideal}})^2 - \sum_{i=1}^m \Delta R_i \cdot \Delta R_i^{\text{ideal}} = 0$$

$$\theta_0^* \cdot \sum_{i=1}^m \Delta R_i^{\text{ideal}} = \sum_{i=1}^m \Delta R_i \cdot \Delta R_i^{\text{ideal}} - \theta_1^* \sum_{i=1}^m (\Delta R_i^{\text{ideal}})^2$$

$$\theta_0^* = \frac{\sum_{i=1}^m \Delta R_i \cdot \Delta R_i^{\text{ideal}} - \theta_1^* \sum_{i=1}^m (\Delta R_i^{\text{ideal}})^2}{\sum_{i=1}^m \Delta R_i^{\text{ideal}}}$$



Model 1:

The parameter θ denotes the slope of the memristors. Negative values represent a negative slope (discordant), values close to 0 are stuck and positive values are more indicative of concordant faults.

Model 2:

In principle similar to model 1 but we now have the intercept. We want this to be close to 0 in order to classify something as ideal.

I would choose model 1 because it is simpler and I am content with the results it delivered.

Memristors were classified with the following code:

```
if(theta > 2.0):
    return MemristorFault.IDEAL
if(theta > 0.1):
    return MemristorFault.CONCORDANT
if(theta < -0.1):
    return MemristorFault.DISCORDANT
else:
    return MemristorFault.STUCK
```

The memristors were classified this way:
Memristor 1 is classified as discordant.
Memristor 2 is classified as concordant.
Memristor 3 is classified as discordant.
Memristor 4 is classified as concordant.
Memristor 5 is classified as stuck.
Memristor 6 is classified as discordant.
Memristor 7 is classified as stuck.
Memristor 8 is classified as concordant.

2 Logistic Regression

The following features were added to the design matrices on top of x1 and x2.

```
center_x = (X_data[:, 0].max() + X_data[:, 0].min()) / 2 #X_data[:, 0] = x1, X_data[:, 1] = x2
center_y = (X_data[:, 1].max() + X_data[:, 1].min()) / 2 #necessary for feature

feature = np.sqrt((X_data[:, 0] - center_x)**2 + (X_data[:, 1] - center_y)**2)
feature1 = X_data[:, 0] * X_data[:, 1]
feature2 = X_data[:, 0] + X_data[:, 1]
```

Features for data set 1

```
X = np.hstack((X_data, X_data ** 2))
```

Features for data set 2

```
sin_x1 = np.sin(x1)
sin_x2 = np.sin(x2)
cos_x1 = np.cos(x1)
cos_x2 = np.cos(x2)
sin_cos_x1 = np.sin(x1) * np.cos(x1)
x1_cubed = x1**3
extra = np.sin(x1)**2 - np.sin(x2)
extra2 = np.sin(x2)**2 - np.sin(x1)
extra3 = np.cos(x1)**2 - np.cos(x2)
extra4 = np.cos(x2)**2 - np.cos(x1)
extra5 = np.sin(x1) + np.cos(x1)
extra6 = np.sin((np.pi * x1) / 2)
```

Features for data set 3

The l2 penalty was used. The following results were achieved:

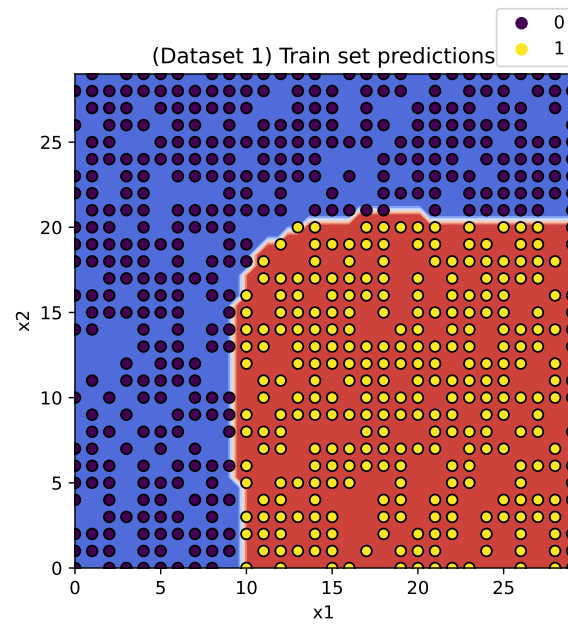

```
---- Logistic regression task 1 ----
Shapes of: X_train (640, 5), X_test (160, 5), y_train (640,), y_test (160,)
Train accuracy: 99.69%. Test accuracy: 98.75%.
Train loss: 0.022944184335746486. Test loss: 0.046497633893654675.
Parameters: [[ 2.70297543 -0.51206517 -0.76460317 -0.21709788  2.19091026]], [-36.32143224]



---- Logistic regression task 2 ----
Shapes of: X_train (640, 4), X_test (160, 4), y_train (640,), y_test (160,)
Train accuracy: 100.00%. Test accuracy: 100.00%.
Train loss: 0.0005398954451803412. Test loss: 0.0038387463242637856.
Parameters: [[ 0.21879489  0.0849633 -0.66633497 -0.66540454]], [394.5933908]

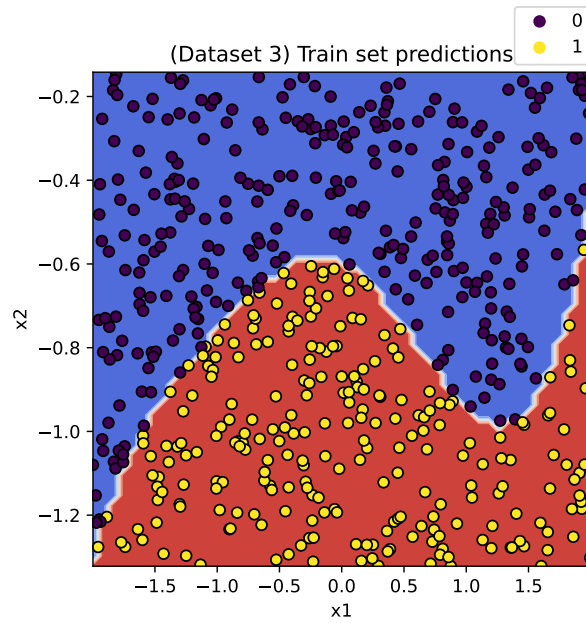
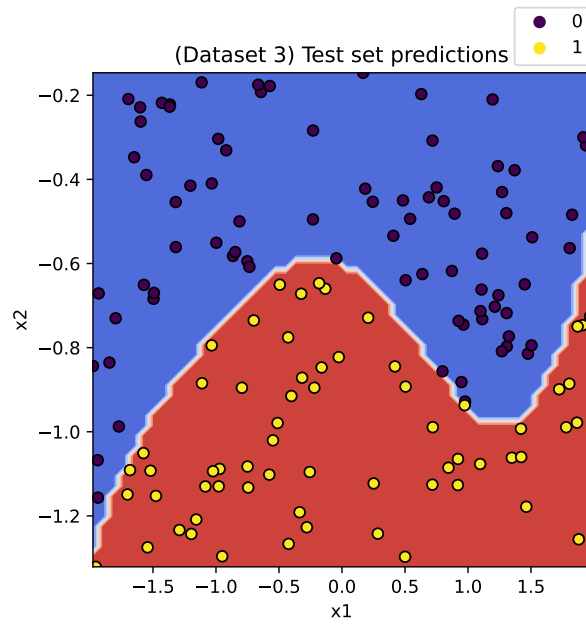
---- Logistic regression task 3 ----
Shapes of: X_train (568, 14), X_test (143, 14), y_train (568,), y_test (143,)
Train accuracy: 92.61%. Test accuracy: 95.10%.
Train loss: 0.17866642615085537. Test loss: 0.1419526673240821.
Parameters: [[-0.5580755 -2.43599138 -0.44461853 -1.94160822 -0.61507153 -1.43984865
 1.9409876 1.28167116 -0.32654179 2.63899903 3.683186 -1.60412163
-1.05969005 0.81839252]], [-1.51649281]
```

Features for data set 3

Resulting plots:







Classifier weights and bias for first data set:

[[2.70297543 -0.51206517 -0.76460317 -0.21709788 2.19091026]], [-36.32143224]

Classifier weights and bias for second data set:

[[0.21879489 0.0849633 -0.66633497 -0.66540454]], [394.5933908]

Classifier weights and bias for third data set:

[[-0.5580755 -2.43599138 -0.44461853 -1.94160822 -0.61507153 -1.43984865
 1.9409876 1.28167116 -0.32654179 2.63899903 3.683186 -1.60412163
 -1.05969005 0.81839252]], [-1.51649281]

Assume we have trained a logistic regression classifier (binary classes) and are given a test dataset D . Is the following statement correct? “If the classifier predicts the correct class for all elements in D (100 percent accuracy), then it follows that the cross-entropy loss (w.r.t. D) is 0.” Explain your reasoning.

This is not necessarily the case because cross-entropy is computed in relation to probability and not the actual classification prediction.

3 Gradient Descent

We set out to minimize the Ackley function. It can be written as:

$$f(x, y) = 20 \cdot e^{-0.2\sqrt{0.5(x^2+y^2)}} - e^{0.5(\cos(2\pi x) + \cos(2\pi y))} + e + 20$$

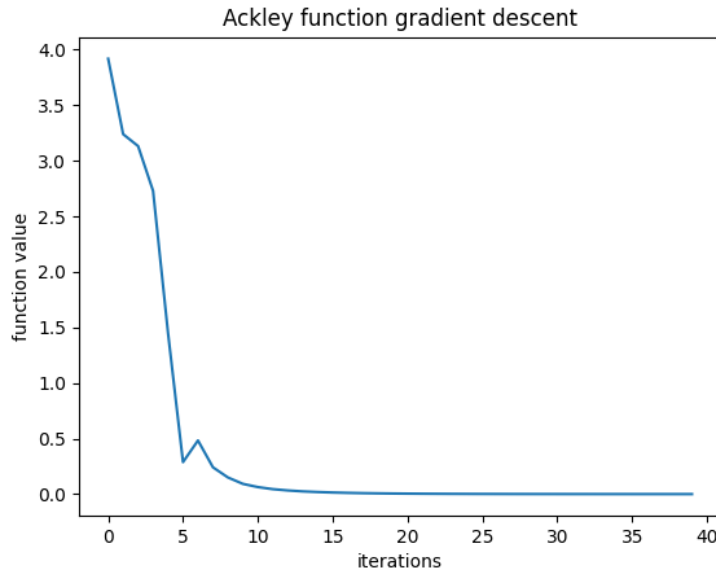
We can reach its global minimum at point (0,0) through gradient descent.

The chosen hyperparameters are as follows:

Number of iterations: 40

Learning rate: 0,1

Learning rate decay: 0,8



The Ackley function is regularly used as a test for optimization algorithms. This function is not convex and it poses a challenge mainly because it has numerous local minima which could "trap" an algorithm and keep it from finding the global minima.

What if we were to use a constant step size?

If we were to set the learning rate decay to 1 and thus have a constant step size it would be very difficult for a function to converge towards the optimum because it would keep "overshooting".

Machine Learning

McVís Omeñudo
Javier Farfán Roca

$$3. f(x,y) = -20 \cdot e^{-0.2\sqrt{0.5(x^2+y^2)}} - e^{0.5(\cos(2\pi x) + \cos(2\pi y))} + e + 20$$

$$\frac{\partial f}{\partial x} = \frac{\partial}{\partial x} -20 \cdot e^{-0.2\sqrt{0.5(x^2+y^2)}} + \frac{\partial}{\partial x} -e^{0.5(\cos(2\pi x) + \cos(2\pi y))} + \frac{\partial}{\partial x} e + \frac{\partial}{\partial x} 20$$

$$\frac{\partial}{\partial x} 20 = 0, \quad \frac{\partial}{\partial x} e = 0$$

$$\frac{\partial}{\partial x} -20 e^{-0.2\sqrt{0.5(x^2+y^2)}} = -20 \cdot (-0.2) \cdot e^{-0.2\sqrt{0.5(x^2+y^2)}} \cdot \frac{1}{2\sqrt{0.5(x^2+y^2)}} \cdot$$

$$\frac{\partial}{\partial x} (0.5(x^2+y^2)) = 4 e^{-0.2\sqrt{0.5(x^2+y^2)}} \cdot \frac{1}{2\sqrt{0.5(x^2+y^2)}} \cdot 0.5 \cdot 2x =$$

$$= 2 e^{-0.2\sqrt{0.5(x^2+y^2)}} \cdot \frac{x}{\sqrt{0.5(x^2+y^2)}}$$

$$\frac{\partial}{\partial x} -e^{0.5(\cos(2\pi x) + \cos(2\pi y))} = -0.5 \cdot e^{0.5(\cos(2\pi x) + \cos(2\pi y))} \cdot \frac{\partial}{\partial x} (\cos(2\pi x))$$

$$= -0.5 \cdot e^{0.5(\cos(2\pi x) + \cos(2\pi y))} \cdot (-2\pi \cdot \sin(2\pi x)) =$$

$$= \pi \cdot e^{0.5(\cos(2\pi x) + \cos(2\pi y))} \cdot \sin(2\pi x)$$

$$\frac{\partial f}{\partial x} = 2 e^{-0.2\sqrt{0.5(x^2+y^2)}} \cdot \frac{x}{\sqrt{0.5(x^2+y^2)}} + \pi \cdot e^{0.5(\cos(2\pi x) + \cos(2\pi y))} \cdot \sin(2\pi x)$$

$$\frac{\partial f}{\partial y} = 2 e^{-0.2\sqrt{0.5(x^2+y^2)}} \cdot \frac{y}{\sqrt{0.5(x^2+y^2)}} + \pi \cdot e^{0.5(\cos(2\pi x) + \cos(2\pi y))} \cdot \sin(2\pi y)$$

Ackley function gradient